

Introduction to R - Day 2

Lidiya Mishieva

29 January, 2026

Intro

- Day 1: basic syntax, classes, objects, functions
- **Day 2: base package, tidy programming & tidyverse**
- Day 3: RMarkdown / Quarto
- Day 4: Git & RStudio

Recap last week

- Assignment
- Object types
- Functions
- Packages

Tidy paradigm - intuitive ‘readable’ syntax

- Consider this story:
 1. *Little bunny Foo Foo*
 2. *Went hopping through the forest*
 3. *Scooping up the field mice*
 4. *And bopping them on the head*

- Translate into code:

- ```
1 # define the object
2 foo_foo <- little_bunny()
3 # define the first operation
4 foo_foo_1 <- hop(foo_foo, through = forest)
5 # define the second operation
6 foo_foo_2 <- scoop(foo_foo_1, up = field_mice)
7 # define the thirs operation
8 foo_foo_3 <- bop(foo_foo_2, on = head)
```

# Tidy paradigm - intuitive ‘readable’ syntax

- `%>%` is the pipe operator (part of `magrittr` package), passing the results of the previous operation to the next:

- ```
1 foo_foo %>%
  2   hop(through = forest) %>%
  3   scoop(up = field_mice) %>%
  4   bop(on = head)
```

- The idea is to do this:

- ```
1 use_this %>%
 2 do_this_first() %>%
 3 do_this_second() %>%
 4 do_this_third()
```

- Instead of this:

- ```
1 do_this_third(do_this_second(do_this_first(use_this)))
```


Tidy paradigm - intuitive ‘readable’ syntax

- R 4.1.0 introduced a *native* pipe operator `|>`
- The essence of the behavior between `%>%` and `|>` is the same
 - “pipe an object forward to a function”
- But there are some differences
 - `%>%` is more complicated, but also more flexible
 - Details about the differences are summarized here:
 - <https://tidyverse.org/blog/2023/04/base-vs-magrittr-pipe/>

Tidy paradigm - tidyverse

Tidyverse

=



R packages for data science

The tidyverse is an opinionated collection of R packages designed for data science. All packages share an underlying design philosophy, grammar, and data structures.

Install the complete tidyverse with:

```
install.packages("tidyverse")
```

Tidy paradigm - pharmaverse

Packages



Tidy paradigm - tibbles and dataframes

- tibbles are very similar to data.frames
- output of tidyverse functions is usually a tibble (if performed on a data.frame)
- unlike data frames, tibbles don't show the entire dataset when you print it:

- ```
1 # load data
2 penguins <- datasets::penguins
3 # object class is data.frame
4 class(penguins)
5 # print the data
6 penguins
7 # convert to tibble
8 penguins <- tibble::as_tibble(penguins)
9 # object class is tibble now
10 class(penguins)
11 # print the data
12 penguins
13 # we can also do the same with dataframes by using head()
14 head(as.data.frame(penguins))
```

# Tidy paradigm - tibbles and dataframes

- tibbles cannot access a column when you provide a partial name of the column, but data frames can

- ```
1 penguins_df <- as.data.frame(penguins)
2 penguins_tibble <- tibble::as_tibble(penguins)
3
4 # works on dataframes
5 penguins_df$sp
6 # does not work on tibbles
7 penguins_tibble$sp
8 penguins_tibble$species
```

Tidy paradigm - tibbles and dataframes

- when you access only one column of a tibble, it will keep the tibble structure. But when you access one column of a data frame, it will become a vector

- ```
1 penguins_tibble[, "species"]
2 class(penguins_tibble[, "species"]) # will stay a tibble
3
4 penguins_df[, "species"]
5 class(penguins_df[, "species"]) # will become a vector
```

# Tidy paradigm - tibbles and dataframes

- when assigning a new column to a tibble, the input will not be recycled, which means you have to provide an input of the same length of the other columns
- but a data frame will recycle the input

```
1 # will return an error
2 penguins_tibble$new_var <- 5:6
3 # will recycle the input across the rows
4 penguins_df$new_var <- 5:6
5 penguins_df$new_var
```

# Tidy paradigm - tibbles and dataframes

- tibbles don't support arithmetic operations on all columns well, the result will be converted into a data frame without any notice

- ```
1 penguins_tibble$new_var <- penguins_tibble$bill_len^2
2 class(penguins_tibble)
```

Note

Take home message: tibbles and data.frames both work well in most cases, just be aware what object class you are working with.

Working with data

Working with data - import/export

- different import functions (and even packages) for different data formats
 - importing csv files:

```
1 # baseR:  
2 read.csv(path)  
3 # tidy:  
4 readr::read_csv(path)  
5 readr::read_csv2(path) # uses a comma as separator
```

- importing excel files:

```
1 readxl::read_excel(path)
```

- importing other common data files :

```
1 haven::read_dta(path)  
2 haven::read_sas(path)  
3 haven::read_sav(path)  
4 haven::read_stata(path)  
5 haven::read_spss(path)
```

Working with data - import/export

- export functions are usually just called `write_` instead of `read_` (check the documentations), but you have to provide the object to be exported and the path in separate arguments, e.g.:

```
1 write_csv(data, file="somepath.csv")
```

- export to excel is not straightforward, better to export csv and then open using excel if needed

Working with data - merging (base)

- **rbind()**
 - stacks rows (adds observations)
 - data frames must have same columns (names & types)
 - example: `rbind(data1, data2)`
- **cbind()**
 - binds columns (adds variables)
 - same number of rows required
 - example: `cbind(data1, new_col)`
- **merge()**
 - join by keys
 - example: `merge(data1, data2, by = "id")`

Working with data - merging (tidy)

- `dplyr::bind_rows()`
 - tidyverse version of `rbind()`
 - handles mismatched columns (fills with `NA`)
 - example: `bind_rows(data1, data2)`
- `dplyr::bind_cols()`
 - tidyverse version of `cbind()`
 - row counts must match
 - example: `bind_cols(data1, data2)`
- **merging (all from dplyr)**
 - `left_join()` & `right_join()`
 - `inner_join()` & `full_join()`

Working with data - long to wide & wide to long (base)

```
1 # load an example dataset from the package datasets (long format)
2 data_long <- datasets::Indometh
3
4 # long to wide
5 data_wide <- reshape(data_long,
6                         idvar = "Subject",
7
8                         timevar = "time",
9                         v.names = "conc",
10
11                         direction = "wide") # direction of reshaping
12
13
14
15
16
17 # wide to long
18 data_long <- reshape(data_wide,
19                         idvar = "Subject",
20
21                         v.names = "conc",
22
23                         direction = "long") # direction of reshaping
```

Working with data - long to wide & wide to long (tidy)

```
1 # load an example dataset from the package datasets (long format)
2 data_long <- datasets::Indometh
3
4 # long to wide
5 data_wide <- data_long %>%
6   pivot_wider(
7     names_from = time,           # which column (or columns) to get the name of the output
8     values_from = conc,         # which column (or columns) to get the cell values from
9     names_prefix = "conc.")    # string added to the start of every variable name
10
11 # wide to long
12 data_long <- data_wide %>%
13   pivot_longer(
14     cols = starts_with("conc"), # columns to pivot into longer format (selection example, use only one)
15     cols = conc.0.25:conc.8,   # columns to pivot into longer format (selection example, use only one)
16
17     names_to = "time",        # a character vector specifying the new column or columns to create
18                           # from the information stored in the column names of data specified by cols
19
20     names_prefix = "conc.",   # a regular expression used to remove matching text from the start of each
21                           # variable name
22
23     values_to = "conc")      # a string specifying the name of the column to create from the data stored
24                           # in cell values
```

Working with data - subsetting/ filtering

- base R
 - `data[rows, cols]`
 - example: `data[data$age > 30, c("name", "age")]`
 - get row indices: `data[which(data$age > 30),]`
- tidy
 - rows: `data %>% filter(age > 30)`
 - columns: `select(data, name, age)`
 - in a pipe: `data %>% select(name, age) %>% filter(age > 30)`
 - helper functions: `starts_with("string")`, `contains("string")`
 - or use regular expressions (regex):
 - <https://r4ds.hadley.nz/regexp.html>

Working with data - renaming variables (base)

- rename by name:
 - `names(data)[names(data) == "old"] <- "new"`
- rename by position:
 - `names(data)[2] <- "new_name"`
- with `colnames()`:
 - `colnames(data)[colnames(data) == "old"] <- "new"`

Working with data - renaming variables (tidy)

- `rename(): data %>% rename(new = old)`
- rename multiple variables at once (`rename_with()` expects a function, indicated by `~`):
 - `data %>% rename_with(~ paste0("var", seq_along(.x)))`
 - `paste0()` # concatenate vectors after converting to character
 - `seq_along(.x)` # generate a sequence along with the length of `.x`
 - `.x` placeholder for the columns of the dataframe; if your dataframe has columns `a, b, c`, then `.x = c("a","b","c")`
 - If we don't use `.x` instead of `seq_along(.x)`, we will concatenate var and variable names from the data instead of var and a sequence number
 - `data %>% rename_with(~ gsub("-", "_", .x))`
 - `gsub()` renames the pattern provided in the first argument, by the argument provided in the second

Working with data - transformations

- base R:
 - add new column:
 - `data$new_var <- data$x * 2`
 - conditional logic:
 - `data$group <- ifelse(data$age > 30, "old", "young")`
- tidy:
 - mutate():
 - `data %>% mutate(new_var = x * 2)`
 - conditional:
 - `data %>% mutate(group = if_else(age > 30, "old", "young"))`
 - multiple variables at once:
 - `data %>% mutate(across(starts_with("x"), log))`

Working with data - modeling

- Linear regression: `lm()`
 - `lm(y ~ x1 + x2, data = data)`
- Some generalized linear models (logistic, poisson): `glm()`
 - `glm(y ~ x1 + x2, data = data, family = binomial)`
 - `glm(y ~ x1 + x2, data = data, family = poisson)`
- Mixed-effects models: `lme4::lmer()` `lme4::glmer()`
 - `lme4::lmer(y ~ x1 + x2 + (1 | group), data = data)`
- Survival models:
 - Cox-model:
 - `survival::coxph(Surv(time, event) ~ x1 + x2, data = data)`
 - Kaplan-Meier:
 - `survival::survfit(Surv(time, event) ~ 1, data = data)`

Next session

- Creating plots and tables
- Reporting with RMarkdown/Quarto

The end



<https://allisonhorst.com/>

Introduction to R - Day 2